

Hi everyone good afternoon!



EE2211 Introduction to Machine Learning

T14 & T22, Chua Dingjuan elechuad@nus.edu.sg
Materials @ tiny.cc/ee2211tut




Tutorial 7

Over/Underfitting, Regularization (Bias, Variance)

Recall... Polynomial Regression

	1	...	d	y	
	ave. hrs of sleep / night			final mark in module	
x_1	5			23	y_1
x_2	4			45	y_2
x_3	3			89	y_3
x_4	8			46	y_4
x_5	9			90	y_5
...	8.5			80	...
x_m	13			30	y_m

final mark in module



In polynomial regression, the system becomes “underdetermined” easily.

Potential for over/under fitting!

$$y = Pw$$

$$\begin{bmatrix} y_1 \\ y_2 \\ \dots \end{bmatrix} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \dots & \dots & \dots \end{bmatrix} \begin{bmatrix} c \\ b \\ a \end{bmatrix} \Rightarrow \begin{bmatrix} 23 \\ 45 \\ \dots \end{bmatrix} = \begin{bmatrix} 1 & 5 & 25 \\ 1 & 4 & 16 \\ 1 & \dots & \dots \end{bmatrix} \begin{bmatrix} c \\ b \\ a \end{bmatrix}$$

Ridge Regression / Regularization


Primal <i>over, λI</i>	Dual Form <i>under, λI</i>
$\hat{\mathbf{w}} = (X^T X + \lambda \mathbf{I})^{-1} X^T \mathbf{y}$	$\hat{\mathbf{w}} = X^T (X X^T + \lambda \mathbf{I})^{-1} \mathbf{y}$

- In ridge regression, an additional term is added during minimization:

$$\min_{\mathbf{w}} \sum_{i=1}^m (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2 + \lambda \mathbf{w}^T \mathbf{w}$$


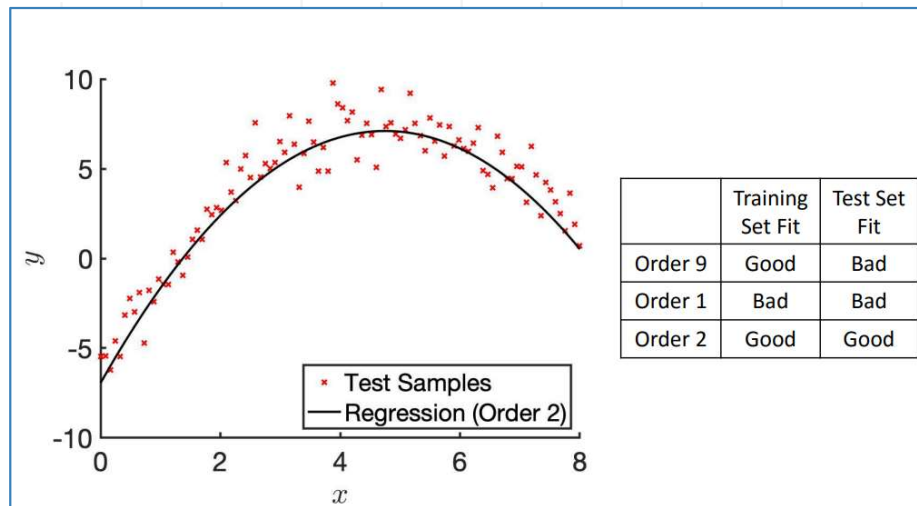
- In applications where number of samples m , much smaller than d (polynomial regression is one example), ridge regression can be useful

- Effect of λ reduces $\mathbf{w} \rightarrow$ same $\Delta x \rightarrow$ smaller Δy
 \rightarrow predictions are less sensitive to training data \rightarrow training error \uparrow




testing error \downarrow





- 1) Underfitting is the inability of trained model to predict the targets in the training set
- 2) Overfitting occurs when model predicts the training data well, but predicts new data (e.g., from test set) poorly -- > (Feature Selection)



Bias-Variance Decomposition Theorem


● Test error = Bias Squared + Variance + Irreducible Noise

● Variance :

Variability of prediction models across different training sets / trained models

● Bias :

How well an average prediction model will perform





Discussion of Solutions

Q1,2

7

Question1

This question explores the use of Pearson's correlation as a feature selection metric. What are the top two features we should select if we use Pearson's correlation as a feature selection metric?

	Datapoint 1	Datapoint 2	Datapoint 3	Datapoint 4	Datapoint 5
Feature 1	0.3510	2.1812	0.2415	-0.1096	0.1544
Feature 2	1.1796	2.1068	1.7753	1.2747	2.0851
Feature 3	-0.9852	1.3766	-1.3244	-0.6316	-0.8320
Target y	0.2758	1.4392	-0.4611	0.6154	1.0006

$$r = \frac{\frac{1}{N} \sum_{n=1}^N (a_i - \bar{a})(b_i - \bar{b})}{\sqrt{\frac{1}{N} \sum_{n=1}^N (a_i - \bar{a})^2} \sqrt{\frac{1}{N} \sum_{n=1}^N (b_i - \bar{b})^2}}$$

Handwritten notes:
 - Above the numerator: covariance
 - Below the first denominator term: std dev of a
 - Below the second denominator term: std dev of b

$$\text{Mean of Feature 1} = \mu_1 = \frac{0.3510 + 2.1812 + 0.2415 - 0.1096 + 0.1544}{5} = 0.5637$$

$$\text{Mean of Feature 2} = \mu_2 = \frac{1.1796 + 2.1068 + 1.7753 + 1.2747 + 2.0851}{5} = 1.6843$$

$$\text{Mean of Feature 3} = \mu_3 = \frac{-0.9852 + 1.3766 - 1.3244 - 0.6316 - 0.8320}{5} = -0.4793$$

$$\text{Mean of Target y} = \mu_y = \frac{0.2758 + 1.4392 - 0.4611 + 0.6154 + 1.0006}{5} = 0.5740$$

$$\text{Feature 1 std} = \sigma_1 = \sqrt{\frac{(0.3510 - \mu_1)^2 + (2.1812 - \mu_1)^2 + (0.2415 - \mu_1)^2 + (-0.1096 - \mu_1)^2 + (0.1544 - \mu_1)^2}{5}} = 0.8229$$

$$\text{Feature 2 std} = \sigma_2 = \sqrt{\frac{(1.1796 - \mu_2)^2 + (2.1068 - \mu_2)^2 + (1.7753 - \mu_2)^2 + (1.2747 - \mu_2)^2 + (2.0851 - \mu_2)^2}{5}} = 0.3924$$

$$\text{Feature 3 std} = \sigma_3 = \sqrt{\frac{(-0.9852 - \mu_3)^2 + (1.3766 - \mu_3)^2 + (-1.3244 - \mu_3)^2 + (-0.6316 - \mu_3)^2 + (-0.8320 - \mu_3)^2}{5}} = 0.9552$$

$$\text{Target y std} = \sigma_y = \sqrt{\frac{(0.2758 - \mu_y)^2 + (1.4392 - \mu_y)^2 + (-0.4611 - \mu_y)^2 + (0.6154 - \mu_y)^2 + (1.0006 - \mu_y)^2}{5}} = 0.6469$$

Question1

This question explores the use of Pearson's correlation as a feature selection metric. What are the top two features we should select if we use Pearson's correlation as a feature selection metric?

	Datapoint 1	Datapoint 2	Datapoint 3	Datapoint 4	Datapoint 5
Feature 1	0.3510	2.1812	0.2415	-0.1096	0.1544
Feature 2	1.1796	2.1068	1.7753	1.2747	2.0851
Feature 3	-0.9852	1.3766	-1.3244	-0.6316	-0.8320
Target y	0.2758	1.4392	-0.4611	0.6154	1.0006

$$r = \frac{\frac{1}{N} \sum_{n=1}^N (a_i - \bar{a})(b_i - \bar{b})}{\sqrt{\frac{1}{N} \sum_{n=1}^N (a_i - \bar{a})^2} \sqrt{\frac{1}{N} \sum_{n=1}^N (b_i - \bar{b})^2}}$$

Handwritten notes:
 - Above the numerator: *covariance*
 - Below the denominator: *std dev of a* and *std dev of b*

Therefore, the top 2 features are Feature 1 and Feature 3.

$$\text{Cov}(\text{Feature 1}, y) = \frac{1}{5}[(0.3510 - \mu_1)(0.2758 - \mu_y) + (2.1812 - \mu_1)(1.4392 - \mu_y) + (0.2415 - \mu_1)(-0.4611 - \mu_y) + (-0.1096 - \mu_1)(0.6154 - \mu_y) + (0.1544 - \mu_1)(1.0006 - \mu_y)] = 0.3188$$

$$\text{Cov}(\text{Feature 2}, y) = \frac{1}{5}[(1.1796 - \mu_2)(0.2758 - \mu_y) + (2.1068 - \mu_2)(1.4392 - \mu_y) + (1.7753 - \mu_2)(-0.4611 - \mu_y) + (1.2747 - \mu_2)(0.6154 - \mu_y) + (2.0851 - \mu_2)(1.0006 - \mu_y)] = 0.1152$$

$$\text{Cov}(\text{Feature 3}, y) = \frac{1}{5}[(-0.9852 - \mu_3)(0.2758 - \mu_y) + (1.3766 - \mu_3)(1.4392 - \mu_y) + (-1.3244 - \mu_3)(-0.4611 - \mu_y) + (-0.6316 - \mu_3)(0.6154 - \mu_y) + (-0.8320 - \mu_3)(1.0006 - \mu_y)] = 0.4949$$

$$\text{Correlation of Feature 1 \& } y = \frac{\text{Cov}(\text{Feature 1}, y)}{\sigma_1 \sigma_y} = \frac{0.3188}{0.8229 \times 0.6469} = 0.5988$$

$$\text{Correlation of Feature 2 \& } y = \frac{\text{Cov}(\text{Feature 2}, y)}{\sigma_2 \sigma_y} = \frac{0.1152}{0.3924 \times 0.6469} = 0.4537$$

$$\text{Correlation of Feature 3 \& } y = \frac{\text{Cov}(\text{Feature 3}, y)}{\sigma_3 \sigma_y} = \frac{0.4949}{0.9552 \times 0.6469} = 0.8009$$

Question1

SOLUTIONS – Possible Python Solutions as suggested by students!

```
#-----  
#EE2211 Tutorial 7 Question 1  
#by Chua Dingjuan  
#-----  
  
import numpy as np  
  
x1 = np.array([0.3510, 2.1812, 0.2415, -0.1096, 0.1544])  
x2 = np.array([1.1796, 2.1068, 1.7753, 1.2747, 2.0851])  
x3 = np.array([-0.9852, 1.3766, -1.3244, -0.6316, -0.8320])  
y=np.array([0.2758, 1.4392, -0.4611, 0.6154, 1.0006])  
  
#corrcoef() returns the correlation matrix, which is a two-dimensional array  
#with the correlation coefficients.  
#here, corrcoef(x1,y) returns correlation of (x1,x1) , then (x1,y), then (y,x1)  
#we only need either (x1,y) / (y,x1) --> r1[0][1]  
r1 = np.corrcoef(x1,y)  
r2 = np.corrcoef(x2,y)  
r3 = np.corrcoef(x3,y)  
  
print('Pearson correlation values are : ', r1[0][1], r2[0][1], r3[0][1])
```

Question2

This question further explores linear regression and ridge regression.

(a) Use the polynomial model from **orders 1 to 6** to train and test the data **without regularization**.

Plot the Mean Squared Errors (MSE) over orders from 1 to 6 for **both the training and the test sets**.

Which model order provides the best MSE in the training and test sets? Why?

(b) Use Regularization $\lambda=1$

SOLUTIONS – Check your answers.

NO REGULARIZATION

Order	1	2	3	4	5	6
Training MSE	2.3071	8.4408e-03	8.3026e-03	1.7348e-03	3.8606e-25	2.3656e-17
Test MSE	3.0006	0.0296	0.0301	0.0854	1.0548	10.7674

REGULARIZATION

Order	1	2	3	4	5	6
Training MSE	2.3586	8.4565e-03	8.3560e-03	1.8080e-03	7.2650e-04	1.9348e-04
Test MSE	3.2756	0.0302	0.0314	0.0939	0.4369	6.0202

Training Data

$\{x = -10\} \rightarrow \{y = 4.18\}$

$\{x = -8\} \rightarrow \{y = 2.42\}$

$\{x = -3\} \rightarrow \{y = 0.22\}$

$\{x = -1\} \rightarrow \{y = 0.12\}$

$\{x = 2\} \rightarrow \{y = 0.25\}$

$\{x = 7\} \rightarrow \{y = 3.09\}$

Testing Data

$\{x = -9\} \rightarrow \{y = 3\}$

$\{x = -7\} \rightarrow \{y = 1.81\}$

$\{x = -5\} \rightarrow \{y = 0.80\}$

$\{x = -4\} \rightarrow \{y = 0.25\}$

$\{x = -2\} \rightarrow \{y = -0.19\}$

$\{x = 1\} \rightarrow \{y = 0.4\}$

$\{x = 4\} \rightarrow \{y = 1.24\}$

$\{x = 5\} \rightarrow \{y = 1.68\}$

$\{x = 6\} \rightarrow \{y = 2.32\}$

$\{x = 9\} \rightarrow \{y = 5.05\}$

Question1

This question further explores linear regression and ridge regression.

(a) Use the polynomial model from **orders 1 to 6** to train and test the data **without regularization**.

Plot the Mean Squared Errors (MSE) over orders from 1 to 6 for **both the training and the test sets**.

Which model order provides the best MSE in the training and test sets? Why?

(a) SOLUTION (Similar to Tutorial 6 Q2(a) just with more orders!)

Training Data

$\{x = -10\} \rightarrow \{y = 4.18\}$

$\{x = -8\} \rightarrow \{y = 2.42\}$

$\{x = -3\} \rightarrow \{y = 0.22\}$

$\{x = -1\} \rightarrow \{y = 0.12\}$

$\{x = 2\} \rightarrow \{y = 0.25\}$

$\{x = 7\} \rightarrow \{y = 3.09\}$

Testing Data

$\{x = -9\} \rightarrow \{y = 3\}$

$\{x = -7\} \rightarrow \{y = 1.81\}$

$\{x = -5\} \rightarrow \{y = 0.80\}$

$\{x = -4\} \rightarrow \{y = 0.25\}$

$\{x = -2\} \rightarrow \{y = -0.19\}$

$\{x = 1\} \rightarrow \{y = 0.4\}$

$\{x = 4\} \rightarrow \{y = 1.24\}$

$\{x = 5\} \rightarrow \{y = 1.68\}$

$\{x = 6\} \rightarrow \{y = 2.32\}$

$\{x = 9\} \rightarrow \{y = 5.05\}$

STEPS :: For each order from 1 to 6

- Create P matrix (polynomial matrix)
- Solve for w (coefficients of regression polynomial, taking note of P's size).

MSE

- For x (training) and x (test), find the corresponding y on your regression line.
- Calculate mean squared errors between that and the y data provided.

Plotting

- Plot training data and testing data as points
- Plot your regression line.

Xtrain (6 data points, 6 rows)

Order	1	2 ..	5	6
P cols	2	3..	6	7
P rows	6	6 ..	6	6

(a) SOLUTION (part 1 - solving)

- 6 training data points
- Polynomial orders 1 to 5 : **Primal** solution $\rightarrow \hat{w} = (P^T P)^{-1} P^T y$
- Order 6 : 7 unknowns, under-determined, so we use the **Dual** solution $\rightarrow \hat{w} = P^T (P P^T)^{-1} y$

STEPS :: For each order from 1 to 6

- Create P matrix (polynomial matrix)
- Solve for w (coefficients of regression polynomial, taking note of P's size).

```
# Create regressors
P_train_list = CreateRegressors(x, max_order)
print(P_train_list)
P_test_list = CreateRegressors(xt, max_order)
P_plot_list = CreateRegressors(x_plot, max_order)
```

```
#####
# Q1 (a)
#####
```

```
# Estimate coefficients WITHOUT REGULARIZATION
w_list = EstimateRegressionCoefficients(P_train_list, y)
```

```
def EstimateRegressionCoefficients(P_list, y, reg=None):
```

```
# P_list is a list
# P_list[i] are regressors for order i+1 and is of size N x (order+1), where
# N is number of data points
```

```
w_list = []
if reg is None:
```

```
    for P in P_list:
        if (P.shape[1] > P.shape[0]): #use dual solution
            w = P.T @ inv(P @ P.T) @ y
        else: # use primal solution
            w = (inv(P.T @ P) @ P.T) @ y
    w_list.append(w)
```

```
else:
```

```
    for P in P_list:
        w = (inv(P.T @ P + reg*np.eye(P.shape[1])) @ P.T) @ y
        w_list.append(w)
```

```
    return w_list
```

```
from sklearn.preprocessing import PolynomialFeatures as skpf
```

```
for order in range(1,7):
```

```
    #step 1, create a polynomial function object class with the right degree
    polyfn = skpf(order)
```

```
    #step 2, fit the polyfn to your actual data and test data.
    P = polyfn.fit_transform(x)
```

```
    #step 3, solving for coefficients of regression polynomial
```

```
    if (P.shape[0] >= P.shape[1]) :
        w = np.linalg.inv(P.T @ P) @ P.T @ y
```

```
    else : #P.shape[0] < P.shape[1], underdetermined systems
        w = P.T @ np.linalg.inv(P @ P.T) @ y
```

```
def CreateRegressors(x, max_order):
```

```
# x is assumed to be array of length N
```

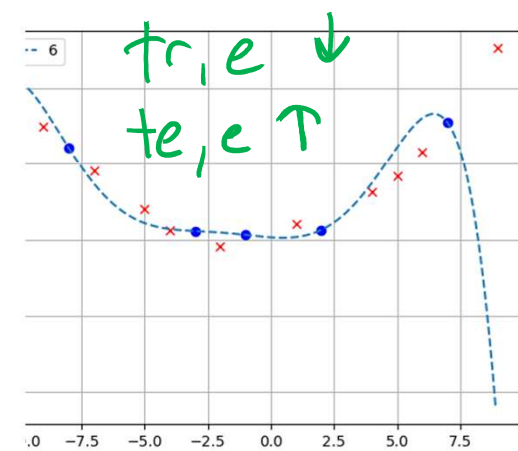

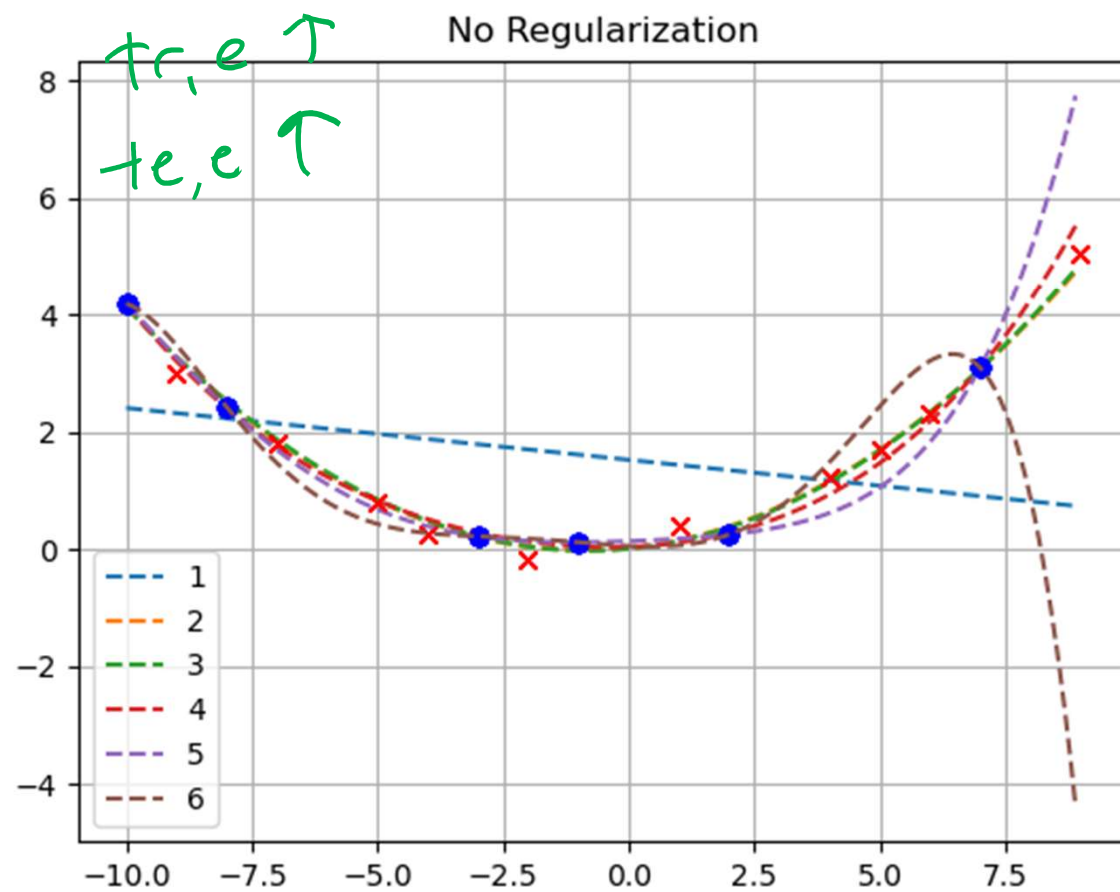
```
# return P = list of regressors based on max_order
```

```
# P[i] are regressors for order i+1 and is of size N x (order+1), where
# N is number of data points
```

```
P = [] #initialize empty list
```

```
for order in range(1, max_order+1):
    current_regressors = np.zeros([len(x), order+1])
    current_regressors[:,0] = np.ones(len(x))
    for i in range(1, order+1):
        current_regressors[:,i] = np.power(x, i)
    P.append(current_regressors)
```

```
return P
```

(a) SOLUTION (part 2 - mse)

MSE

- For x (training) and x (test), find the corresponding y on your regression line.
- Calculate mean squared errors between that and the y data provided.

```
# Apply prediction: predictions are N x max_order
y_train_pred = PerformPrediction(P_train_list, w_list)
y_test_pred  = PerformPrediction(P_test_list, w_list)
y_plot_pred  = PerformPrediction(P_plot_list, w_list)
```

```
#step 4, using regression coefficients to calculate predicted values of y
Ptest = polyfn.fit_transform(xtest)
```

```
y_reg = P @ w
ytest_reg = Ptest @ w
```

```
#step 5, calculating mse error for this order
#from sklearn.metrics import mean_squared_error as skmse
```

```
trainingmse.append(skmse(y, y_reg))
testmse.append(skmse(ytest, ytest_reg))
```

```
def PerformPrediction(P_list, w_list):
```

```
# P_list is list of regressors
# w_list is list of coefficients
# Output is y_predict_mat which N x max_order, where N is the number of samples
```

```
N = P_list[0].shape[0]
max_order = len(P_list)
y_predict_mat = np.zeros([N, max_order])
for order in range(len(w_list)):
    y_predict = np.matmul(P_list[order], w_list[order])
    y_predict_mat[:, order] = y_predict

return y_predict_mat
```

```
# Compute MSE
```

```
train_error = y_train_pred - np.matlib.repmat(y, max_order, 1).T
train_MSE   = np.power(train_error, 2)
train_MSE   = np.mean(train_MSE, 0)
```

```
test_error = y_test_pred - np.matlib.repmat(yt, max_order, 1).T
test_MSE   = np.power(test_error, 2)
test_MSE   = np.mean(test_MSE, 0)
```

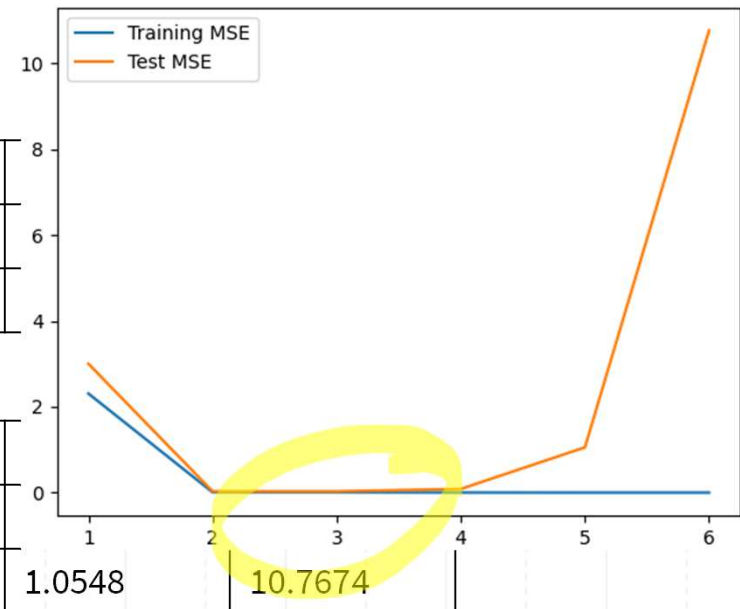
(a) SOLUTION (part 2 - mse)

Official Solutions


Order	1	2	3	4
Training MSE	2.3071	8.4408e-03	8.3026e-03	1.7348e-03
Test MSE	3.0006	0.0296	0.0301	0.0854

IDLE Compiler Produces the following for same code.

Order	1	2	3	4
Training MSE	2.3071	8.4408e-03	8.3026e-03	1.7348e-03
Test MSE	3.0006	0.0296	0.0301	0.0854



- Very low training MSE for orders 5 & 6 but high test MSE → Overfitting
- High training and test MSE for order 1 → Underfitting
- Orders 2 to 4 produce relatively low MSEs (even though model itself is order 2)



(a) SOLUTION (part 3 - plotting)

```
#--- Plotting---#  
plt.figure(1)  
plt.plot(x,y,'bo')  
plt.plot(xtest,ytest,'rx')
```


Training data
Test data

```
xline= np.arange(min(np.concatenate((x,xtest))),max(np.concatenate((x,xtest))),  
xline_resaped = xline.reshape((len(xline),1))  
Pline= polyfn.fit_transform(xline_resaped)  
yline=Pline @ w  
plt.plot(xline,yline,'--', label=order)  
plt.legend()  
plt.grid()
```

Regression Line
P Matrix
Dotted Line

```
plt.figure(2)  
plt.plot(np.arange(1,7),trainingmse,label="Training MSE")  
plt.plot(np.arange(1,7),testmse,label="Test MSE")  
plt.legend()  
plt.show()
```

Training Errors



Question2 (b)

(b) Use Regularization (ridge regression) $\lambda=1$ for all orders and repeat the same analyses. Compare the plots of (a) and (b). What do you see?

(b) SOLUTION

- With regularization, we can simply use the primal solution even for order 6
- Only one small difference in code:

```
#step 3, solving for coefficients of regression polynomial  
w = np.linalg.inv(P.T @ P + np.identity(len(P.T))) @ P.T @ y
```

NO REGULARIZATION

Order	1	2	3	4	5	6
Training MSE	2.3071	8.4408e-03	8.3026e-03	1.7348e-03	3.8606e-25	2.3656e-17
Test MSE	3.0006	0.0296	0.0301	0.0854	1.0548	10.7674

REGULARIZATION

Order	1	2	3	4	5	6
Training MSE	2.3586	8.4565e-03	8.3560e-03	1.8080e-03	7.2650e-04	1.9348e-04
Test MSE	3.2756	0.0302	0.0314	0.0939	0.4369	6.0202

- None of the curves pass through training data exactly → training error increased slightly
- Test MSE for orders 5 & 6 reduced → Overfitting reduced.
- But Test MSE for orders 1 to 4 increased!! → Overly strong regularization.
- Regularization did not help in identifying best polynomial order → cross-validation to be learnt in future!

